

Nishad Mansoor

732-781-6631 | nishadmansoor02@gmail.com | [linkedin.com/nishad-mansoor](https://www.linkedin.com/in/nishad-mansoor) | github.com/nishadmansoor

Education

Northeastern University

MS in Data Science

Boston, MA

Expected December 2026

Rutgers University

BS in Finance and Business Analytics/Information Technology

New Brunswick, NJ

May 2024

Experience

AI Engineer Intern

June 2026 - August 2026

EisnerAmper

Iselin, NJ

- Built a pipeline tracking AI usage across the firm via the Anthropic and OpenAI APIs, covering 2,400 employees
- Developed a time-series forecasting model to project \$1.7M in year-end AI spend, correcting an initial linear model that failed to capture adoption saturation
- Implemented a conversational agent that answers usage/cost questions via natural language by querying a Power BI semantic model hosted on Azure using DAX queries

Machine Learning Researcher

January 2026 – June 2026

Massachusetts Institute of Technology

Cambridge, MA

- Benchmarked 4 VLMs across 300+ origami models using IoU and CLIP similarity, surfacing spatial reasoning patterns
- Achieved query efficiency scores of 0.85–0.87 by fine-tuning constraint thresholds on model outputs
- Co-authored OrigamiBench, accepted to ICML 2026 LM4Plan and CVPR 2026 MAR Workshops

Data Science Intern

October 2025 – May 2026

Defense Industrial Base Institute at Northeastern University

Boston, MA

- Built an ingestion pipeline to pull DoD contract data from USASpending.gov and geocode vendor locations
- Engineered a Neo4j knowledge graph connecting vendors, contracts, and geographies via UEI/CAGE identifiers
- Transformed 120+ MILCON and OSD PDFs into structured datasets via Python and SQL, cut manual review by 60%.

Graduate Research Assistant

January 2025 – May 2026

Northeastern University

Boston, MA

- Engineered Python pipelines for HTML document parsing and CSV generation, reduced prep time by 70%
- Created an AI powered research dashboard using Python and Claude, improved research accessibility by 40%

Data Science Intern

November 2024 - April 2025

Beats Electronics

Remote

- Trained a Random Forest model on marketing spend and 1,000 surveys, achieved F1-score of 0.84 for audience targeting
- Developed a TF-IDF NLP sentiment pipeline on 5,000 reviews, identifying customer satisfaction drivers

Projects

Easy Code | Python, Qdrant, Neo4j, AST, RAG | [GitHub](#)

June 2026 - Present

- Engineering a hybrid code retrieval system that parses GitHub repositories with AST and indexes code in Qdrant for semantic search
- Building a Neo4j knowledge graph linking files, classes, functions, methods and imports to map codebase structure
- Developing REST API combining vector and graph search to return relevant code with file and line references

Reverse Turing Test | Python, PyTorch, DistilBERT, Flask, Docker | [GitHub](#)

February 2026 – May 2026

- Built a 3 model AI-authorship detection framework (DistilBERT: 98.8%, stylometric: 76.8%, visual CNN: 97.5%) on 156k samples from the RAID dataset
- Designed a leakage-free protocol — base models on train, meta-classifier on validation; benchmarked every model against a 51.5% majority-class baseline
- Deployed a containerized Flask app to Hugging Face Spaces with real-time inference, prediction-flow visualization, and a PCA feature-space explorer

Technical Skills

Languages: Python, SQL, R

ML/NLP: PyTorch, Hugging Face, scikit-learn, OpenCV, Pandas, Numpy

Tools: Neo4j, Qdrant, Docker, Git, Streamlit, Flask, Azure